

Monte Carlo ES for Blackjack

Jingyi Zhao

October 2022

1 Problem Setup: State space and Action space

Let S denote the state space and A denote the action space.

Let $s = (i, j) \in S$ where i represents the player's sum and j represents the dealer's sum.

Let $a \in A = \{0, 1\}$ where $a = 1$ indicates *hit* and $a = 0$ indicates *stick*.

2 Pseudocode for Main Algorithms

Algorithm 1 Monte Carlo ES (Exploring Starts) for Blackjack

Initialize:

Create $\pi, Q, Returns$ to be empty dictionaries

$\pi[s] = 1$ if $\text{player_sum} \geq 20$ otherwise 0, for all $s \in S$

$Q[(s, a)] = 0$, for all $s \in S, a \in A$

$Returns[(s, a)] = []$ (empty list), for all $s \in S, a \in A$

loop for NUM_LOOP of times:

Choose $S_0 \in S, A_0 \in A$ randomly such that all pairs have possibility > 0

Generate an episode from S_0, A_0 and π : $S_0, A_0, S_1, A_1, \dots, S_{T-1}, A_{T-1}, R_T$

via Algorithm2

$G \leftarrow R_T$

for $t = T-1, T-2, \dots, 0$ **do**

Append G to $Returns[(S_t, A_t)]$

$Q[(S_t, A_t)] \leftarrow \text{average}(Returns[(S_t, A_t)])$

if $Q[(S_t, 0)] > Q[(S_t, 1)]$ **then**

$\pi[S_t] = Q[(S_t, 0)]$

else

if $Q[(S_t, 0)] < Q[(S_t, 1)]$ **then**

$\pi[S_t] = Q[(S_t, 1)]$

end if

end if

end for

end loop

Algorithm 2 Generate Episode given S_0, A_0, π

Initialize:

$s = S_0, a = A_0, episode = [S_0, A_0]$

while $a = 1$ **do**

$player_sum \leftarrow player_sum + \text{a random card value}$

if $player_sum$ bursts **then**

$episode.append(-1)$

return $episode$

end if

 Update s

$a \leftarrow \pi[s]$

 Append s, a to $episode$

end while

while $dealer_sum \leq 16$ **do**

$dealer_sum \leftarrow dealer_sum + \text{a random card value}$

if $dealer_sum$ bursts **then**

$episode.append(1)$

return $episode$

end if

end while

Compare $player_sum$ with $dealer_sum$, and append return to $episode$ accordingly

return $episode$

3 Experiment Results

3.1 Case 1

After training $n \times 1000$ steps, run 100 episodes with the current policy (no random initial action) and calculate the average return for those 100 episodes.

Provide the average returns for $n = 1, 2, \dots, 20$.

Figure 1 shows the average returns under different n 's.

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
average return	0.91	0.94	-0.17	-0.36	0.15	0.05	0.84	0.59	0.58	0.31	-0.25	-0.49	-0.33	-0.48	0.14	-0.46	-0.61	0.03	0.33	-0.12

Figure 1: Returns for Case 1

Figure 2 shows the optimal policy under this training.

Optimal Policy




Figure 2: Optimal Policy for Case 1

Recall that 1 indicates *hit* and 0 indicates *stick*.

The horizontal direction shows the dealer's showing, where the index 0 here means dealer's showing is 2, the index 1 here means dealer's showing is 3, ..., the index 9 here means dealer's showing is 11.

The vertical direction shows the player's sum, where the index 0 here means player's sum is 12, the index 1 here means the player's sum is 13, ..., the index 10 here means the player's sum is 22.

For instance, the grid in the upper left corner indicates that under our policy, we choose *hit* when player's sum is 12 and dealer's showing is 2.

3.2 Case 2

- Train for 100,000 steps (instead of 20,000 steps)

- Every 1000 steps, evaluate the current policy by running 1000 episodes and averaging over the 1000 episodes

Figure 3 shows the average returns.




Figure 3: Returns for Case 2

Figure 4 shows the optimal policy under this training.

Optimal Policy




Figure 4: Optimal Policy for Case 2